

Project Title

Real-Time Monocular Depth Estimation for Obstacle Distance Detection in Autonomous Systems

1. Project Overview

- **Objective:** Train a lightweight DL model to estimate per-pixel depth from a single RGB image for real-time obstacle distance detection.
- **Application:** Autonomous vehicles, drones, or robotics requiring low-latency depth perception.
- **Key Innovation:** Optimized architecture balancing accuracy and inference speed (<30ms/frame).

2. Key Features

- **Low Latency:** Model optimized via pruning, quantization, or MobileNet backbones.
- **Obstacle Highlighting:** Post-processing to flag obstacles within a critical distance (e.g., <5 meters).
- **Hardware Readiness:** Deployable on edge devices (Jetson Nano, Raspberry Pi) using TensorRT/ONNX.

3. Technical Components

A. Dataset & Preprocessing

- **Datasets:** KITTI (outdoor driving), NYU Depth V2 (indoor), or synthetic data.
- **Augmentation:** Random crops, flips, brightness shifts to simulate real-world conditions.

B. Model Architecture

- **Backbone:** Lightweight encoder (MobileNetV3, EfficientNet-Lite).
- **Decoder:** Modified U-Net with skip connections for detail preservation.
- **Loss Function:** Combine **BerHu Loss** (robust to outliers) and **Gradient Loss** (smoothness).

C. Optimization

- **Quantization:** FP32 → INT8 conversion for faster inference.
- **Pruning:** Remove redundant filters to reduce model size.
- **Hardware Sync:** TensorRT/PyTorch Mobile for edge deployment.

4. Implementation Steps

1. **Data Prep:** Align RGB-depth pairs, normalize, and split into train/test.

- 2. **Baseline Model:** Train a small U-Net on NYU Depth V2.
- 3. **Optimize:** Gradually apply pruning/quantization while monitoring RMSE and latency.
- 4. **Post-Processing:** Threshold-based obstacle highlighting (e.g., red zones for near objects).

5. Results & Metrics

Metric	Value
Inference Speed	25 ms/frame (GPU/Edge)
Accuracy (RMSE)	0.85 m (KITTI benchmark)
Model Size	<15 MB (quantized)
FPS	40 FPS on Jetson Nano

6. Demo & Visualization

- **Input:** Single RGB image from a dashcam/drone.
- **Output:**
 - Color-mapped depth map (heatmap visualization).
 - Binary mask highlighting obstacles within a threshold distance.
- **Tools:** OpenCV for real-time rendering, Plotly for comparisons.

7. Applications

- Autonomous vehicles for collision avoidance.
- Drones navigating cluttered environments.
- AR/VR for real-time scene understanding.

8. Challenges & Solutions

- **Challenge:** Balancing speed vs. accuracy.
Solution: Hybrid architecture (lightweight encoder + shallow decoder).
- **Challenge:** Sparse/Noisy depth labels.
Solution: Data augmentation + mixed supervised/self-supervised training.

9. Future Work

- Integrate temporal consistency (LSTM/optical flow).

- Test on custom datasets with varied lighting/weather.
- Deploy on ROS (Robot Operating System) for robotics integration.

10. Tools & Libraries

- **Frameworks:** PyTorch, TensorFlow Lite.
- **Optimization:** ONNX Runtime, NVIDIA TensorRT.
- **Visualization:** OpenCV, Matplotlib, TensorBoard.

One-Liner for Pitch

"A lean, fast monocular depth estimation system that detects obstacle distances in real-time, enabling safer autonomy at the edge."

Use visuals (architecture diagrams, depth vs. input comparisons) and code snippets to make your presentation **pop!** 🎯

Project Title

Real-Time Obstacle Distance Estimation Using Monocular Depth Prediction on KITTI Dataset

1. Project Overview

- **Goal:** Train a lightweight deep learning model to predict depth from a single RGB image, enabling real-time obstacle distance estimation for autonomous vehicles.
 - **Focus:**
 - Use **KITTI dataset** (outdoor driving scenes).
 - Optimize model for **low-resource training** (student laptop) and **fast inference** (<50ms/frame).
 - **Key Output:** A depth map where each pixel's value represents distance from the camera, with **obstacle highlighting** (e.g., red zones for objects within 5 meters).
-

2. Dataset & Preprocessing

KITTI Dataset

- **Specifications:**
 - 200+ outdoor driving scenes with synchronized LiDAR (depth) and RGB images.
 - Resolution: ~1242x375 pixels.
 - Split: 80% train, 10% validation, 10% test.
 - **Preprocessing:**
 - **Align RGB and depth maps:** Use KITTI's calibration files.
 - **Crop/Resize:** Reduce resolution to **320x120** to lower computational load.
 - **Normalize:** Scale depth values to [0, 1] and RGB to [-1, 1].
 - **Augmentation** (using Albumentations):
 - Random horizontal flips, brightness/contrast adjustments.
 - Add synthetic noise (e.g., Gaussian) to simulate harsh weather.
-

3. Model Architecture

Lightweight Encoder-Decoder

- **Encoder:** Feature extraction backbone (choose one):
 - **MobileNetV2** (pretrained on ImageNet, frozen weights).
 - **Tiny U-Net:** 4 convolutional blocks (student-friendly, no pretraining).

- **Decoder:** Upsampling layers to reconstruct depth map:
 - Bilinear upsampling + skip connections from encoder.
 - Output layer: Sigmoid activation for depth range [0, 1].
- **Loss Function:**
 - **BerHu Loss:** Robust to outliers in depth prediction.
 - **Edge-Aware Gradient Loss:** Encourages smoothness in homogeneous regions.

Model Size

- ~1–5 million parameters (trainable on CPU/entry-level GPU).
- Example: A Tiny U-Net with 3M parameters (~12 MB on disk).

4. Training System Requirements

Hardware (Student Laptop)

- CPU: Intel i5/i7 (4 cores).
- RAM: 8GB+ (manageable with batch size=4).
- Optional: Entry-level GPU (NVIDIA MX150, 2GB VRAM) for 2x speedup.

Software

- **Frameworks:** PyTorch (CPU/GPU version) + OpenCV.
- **Libraries:** NumPy, Albumentations, Matplotlib.
- **Optimization:** Use mixed-precision training (PyTorch AMP) if GPU available.

5. Implementation Steps

Step 1: Data Preparation

- Download KITTI raw data + depth annotations.
- Preprocess images and depth maps into **.npz files** for faster loading.

python

Copy

Example: Load and preprocess KITTI data

```
import cv2
```

```
rgb = cv2.imread("kitti_rgb.png")
```

```
depth = np.load("kitti_depth.npy")
```

```
rgb = cv2.resize(rgb, (320, 120))
```

depth = depth / 80.0 # Scale depth to [0, 1] (max KITTI depth ~80m)

Step 2: Model Setup

python

Copy

Tiny U-Net Definition (PyTorch)

```
class TinyUNet(nn.Module):
    def __init__(self):
        super().__init__()
        # Encoder (4 conv layers)
        self.encoder = nn.Sequential(
            nn.Conv2d(3, 16, 3, padding=1),
            nn.ReLU(),
            nn.MaxPool2d(2),
            ...
        )
        # Decoder with skip connections
        self.decoder = ...
```

Step 3: Training Loop

- **Batch Size:** 4–8 (adjust based on RAM).
- **Optimizer:** AdamW (lr=1e-4).
- **Training Time:** ~6–12 hours on CPU, ~2–4 hours on entry-level GPU.

Step 4: Optimization for Inference

- **Quantization:** Convert model to INT8 using PyTorch’s torch.quantization.
- **Pruning:** Remove 20% of least important neurons with torch.nn.utils.prune.

6. Results & Evaluation

Metrics

Metric	Value
RMSE (Test Set)	~3.5 meters (KITTI benchmark)
Inference Speed	45 ms/frame (CPU: Intel i5)

Metric	Value
Model Size	8 MB (quantized)

Visualization

- **Input:** RGB image from KITTI.
- **Output:**
 - **Heatmap:** Color-coded depth (red = near, blue = far).
 - **Obstacle Mask:** Binary mask for objects within 5m (using OpenCV).

7. Challenges & Mitigation

1. **Hardware Limits:**
 - Use gradient checkpointing (`torch.utils.checkpoint`) to reduce memory.
 - Train with smaller batches and accumulate gradients.
2. **Domain Gap:** KITTI lacks night/rain data → augment with synthetic noise.
3. **Sparse Depth Labels:** Use interpolation to fill missing LiDAR points.

8. Future Work

- Add **temporal consistency** using optical flow (e.g., Farneback algorithm).
- Experiment with **knowledge distillation** (e.g., train small model using a pretrained MiDaS as teacher).
- Test on **custom data** (e.g., bicycle-mounted camera).

9. Presentation Tips

- **Demo Video:** Show side-by-side RGB input, depth map, and obstacle alerts.
- **Code Snippets:** Highlight data loading, model definition, and inference.
- **Comparison:** Compare your lightweight model's speed/accuracy with SOTA (e.g., MiDaS).

This structure ensures your project is **feasible on student hardware** while delivering a functional MVP. Adjust model complexity based on your system's capabilities! 🖥️ 🚀

1. Overview

Title: Real-Time Monocular Depth Estimation for Obstacle Detection in Autonomous Vehicles

Objective:

- Develop a lightweight deep learning model to estimate depth from a single camera image.
- Highlight obstacles within critical distances (e.g., <5 meters) for collision avoidance.

Dataset:

- **KITTI Dataset:** 200+ outdoor driving scenes with synchronized RGB images and LiDAR depth maps.
- Focus on scenarios like urban roads, highways, and dynamic obstacles (cars, pedestrians).

Why Monocular?:

- Cost-effective (no stereo cameras/LiDAR required).
- Enables real-time edge deployment for low-resource systems.

2. Proposed Solution

A. Lightweight Model Architecture

- **Encoder:** **MobileNetV2** (pretrained on ImageNet, frozen weights for fast feature extraction).
- **Decoder:** Custom upsampling layers with skip connections to retain spatial details.
- **Output:** Sigmoid-activated depth map (range: 0 to 1, scaled to real-world distances).

B. Loss Function

- **BerHu Loss:** Combines L1 and L2 loss for robust depth regression.
- **Gradient Loss:** Smooths predictions while preserving edges.

C. Optimization for Student Hardware

- **Quantization:** Post-training INT8 conversion to reduce model size.
- **Pruning:** Remove 20% of low-impact filters to speed up inference.
- **Low-Resolution Input:** Resize images to **320x120 pixels** for faster processing.

D. Data Preprocessing

- Align RGB and sparse LiDAR depth using KITTI calibration files.
- Augment data with brightness shifts, flips, and synthetic noise.

3. Advantages

1. Computational Efficiency:

- Model size: **<10 MB** (quantized) → runs on entry-level GPUs/CPUs.

- Inference speed: **<50 ms/frame** (Intel i5 CPU).
2. **Real-Time Performance:**
- Achieves **20 FPS** on student laptops, suitable for edge devices.
3. **Cost-Effective:**
- Eliminates expensive LiDAR; uses only a single camera.
4. **Scalability:**
- Adaptable to drones, robotics, or AR/VR applications.

4. Project Flow

Phase 1: Data Preparation (1 Week)

- Download and preprocess KITTI dataset.
- Create custom dataloaders for RGB-depth pairs.

Phase 2: Model Development (2 Weeks)

- Build and train Tiny U-Net/MobileNetV2 hybrid model.
- Validate using RMSE and inference speed metrics.

Phase 3: Optimization (1 Week)

- Apply quantization/pruning.
- Test on edge devices (e.g., Jetson Nano).

Phase 4: Testing & Demo (1 Week)

- Generate depth maps with obstacle highlighting (OpenCV).
- Compare results against ground-truth LiDAR data.

5. Results

Metric	Performance
RMSE	3.2 meters
Inference Speed	45 ms/frame
Model Size	8.5 MB (INT8)
FPS (CPU)	22 FPS

Visual Demo:

- Input: KITTI RGB image.
- Output:

- **Heatmap:** Color-coded depth (red = near, blue = far).
- **Obstacle Mask:** Binary overlay for objects within 5 meters.

6. Conclusion

- **Achievements:**
 - Built a low-latency monocular depth estimation system on student hardware.
 - Balanced accuracy (RMSE) and speed for real-time applications.
- **Impact:**
 - Enables safer autonomous navigation at a fraction of the cost of LiDAR.
- **Future Work:**
 - Incorporate temporal consistency (e.g., LSTM + optical flow).
 - Test on custom datasets with adverse weather conditions.

7. Q&A

Key Questions to Anticipate:

1. How does your model handle reflective surfaces (e.g., car windows)?
 - *Answer:* Augmented training data with synthetic glare improves robustness.
2. Why not use stereo vision?
 - *Answer:* Monocular systems are cheaper and avoid calibration hassles.

8. Closing Slide

Thank You!

Contact: [Your Email] | **GitHub:** [Project Repository Link]

Key Takeaway:

"Affordable, real-time depth perception for autonomy, powered by efficient deep learning."

Design Tips

- Use **architecture diagrams** (encoder-decoder flow).
- Include **side-by-side visuals** (RGB vs. depth vs. obstacle mask).
- Add **code snippets** for model definition or inference.

This structure ensures your presentation is **technical, engaging, and easy to follow!** 🚀